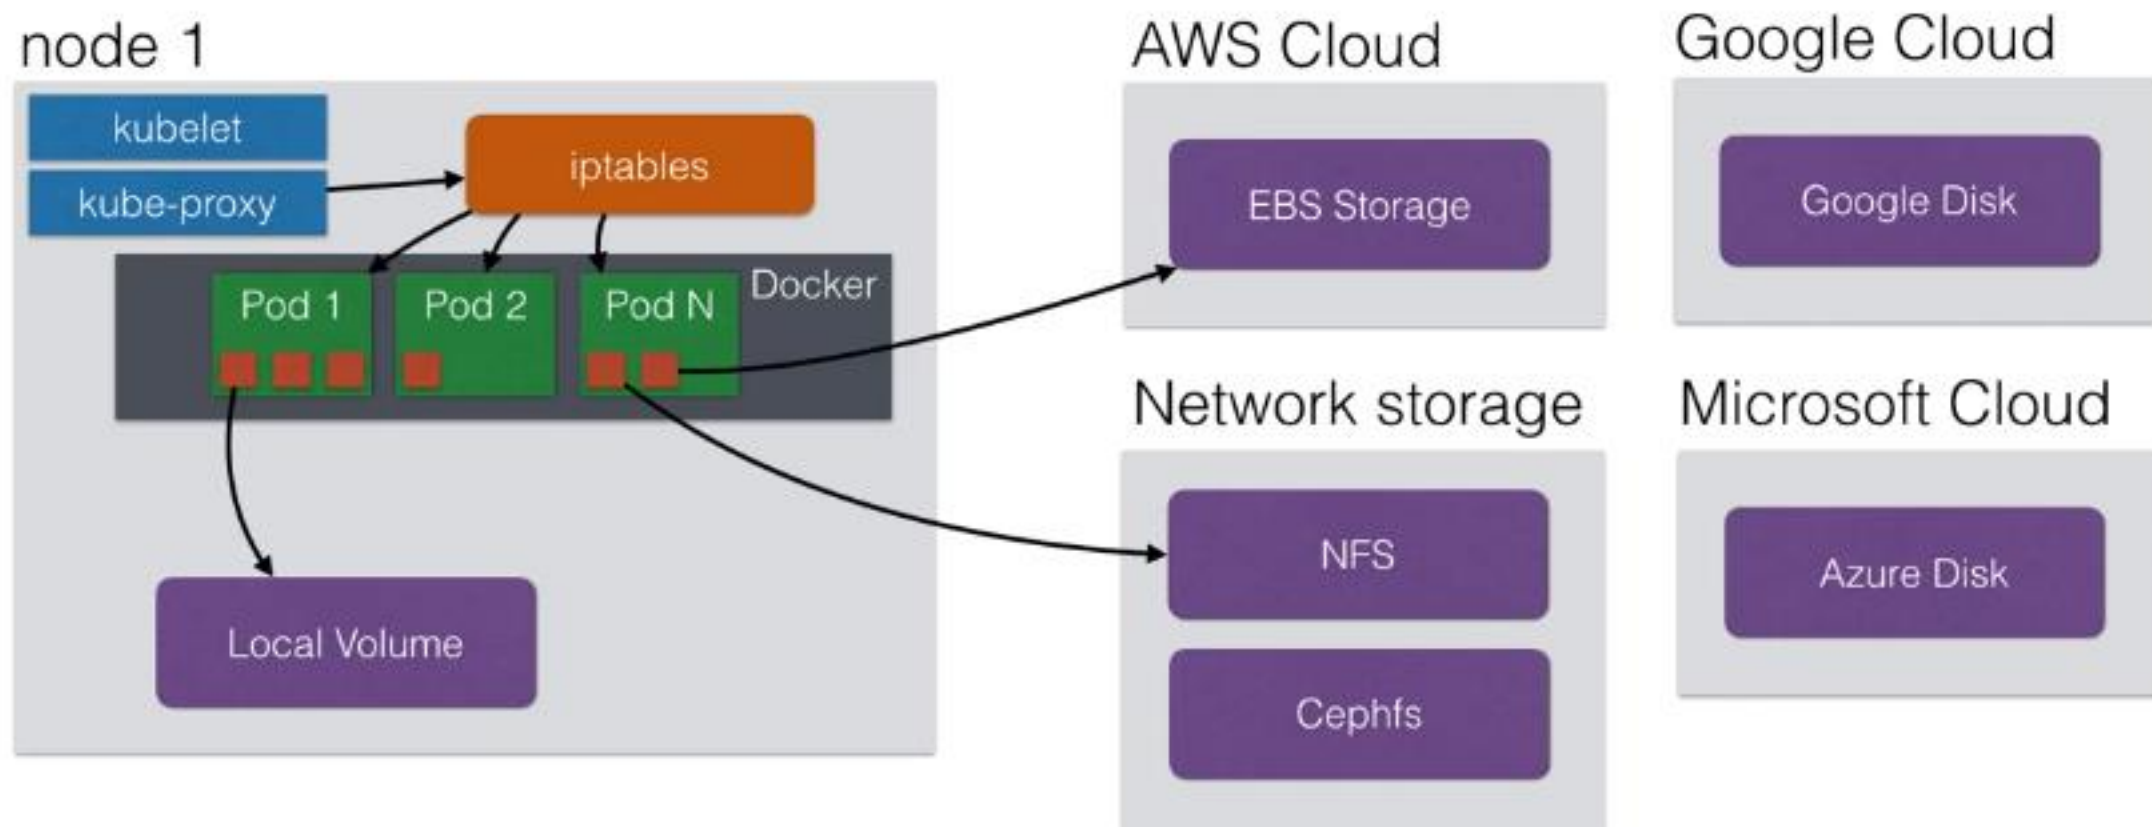


Volumes

- Volumes in kubernetes allow you to **store data outside the container**
- When a containers **stops**, all data on the container itself is **lost**
 - That's why up until now I've been using **stateless** apps: apps that don't keep a **local** state, but store theirs state in an **external service**
 - External Service like a database, caching server (e.g. MySQL, AWS S3)
- Persistent Volumes in Kubernetes allow you **attach a volume** to a container that will **exists** even when the **container** stops

Kubernetes Volumes

- Volumes can be attached using different volumes plugins:

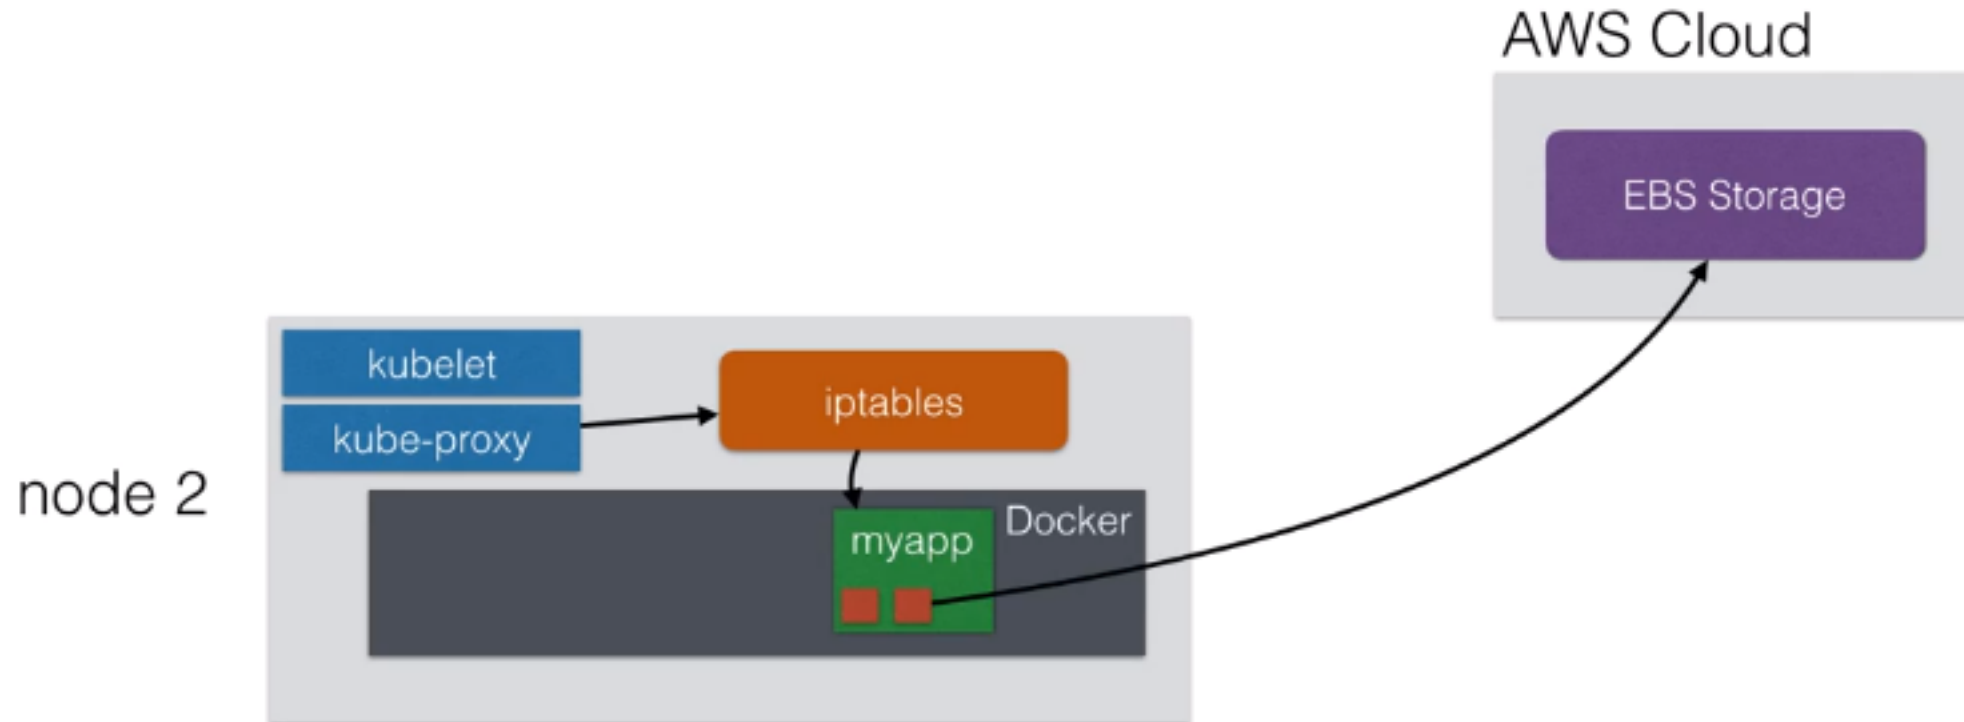


Volumes

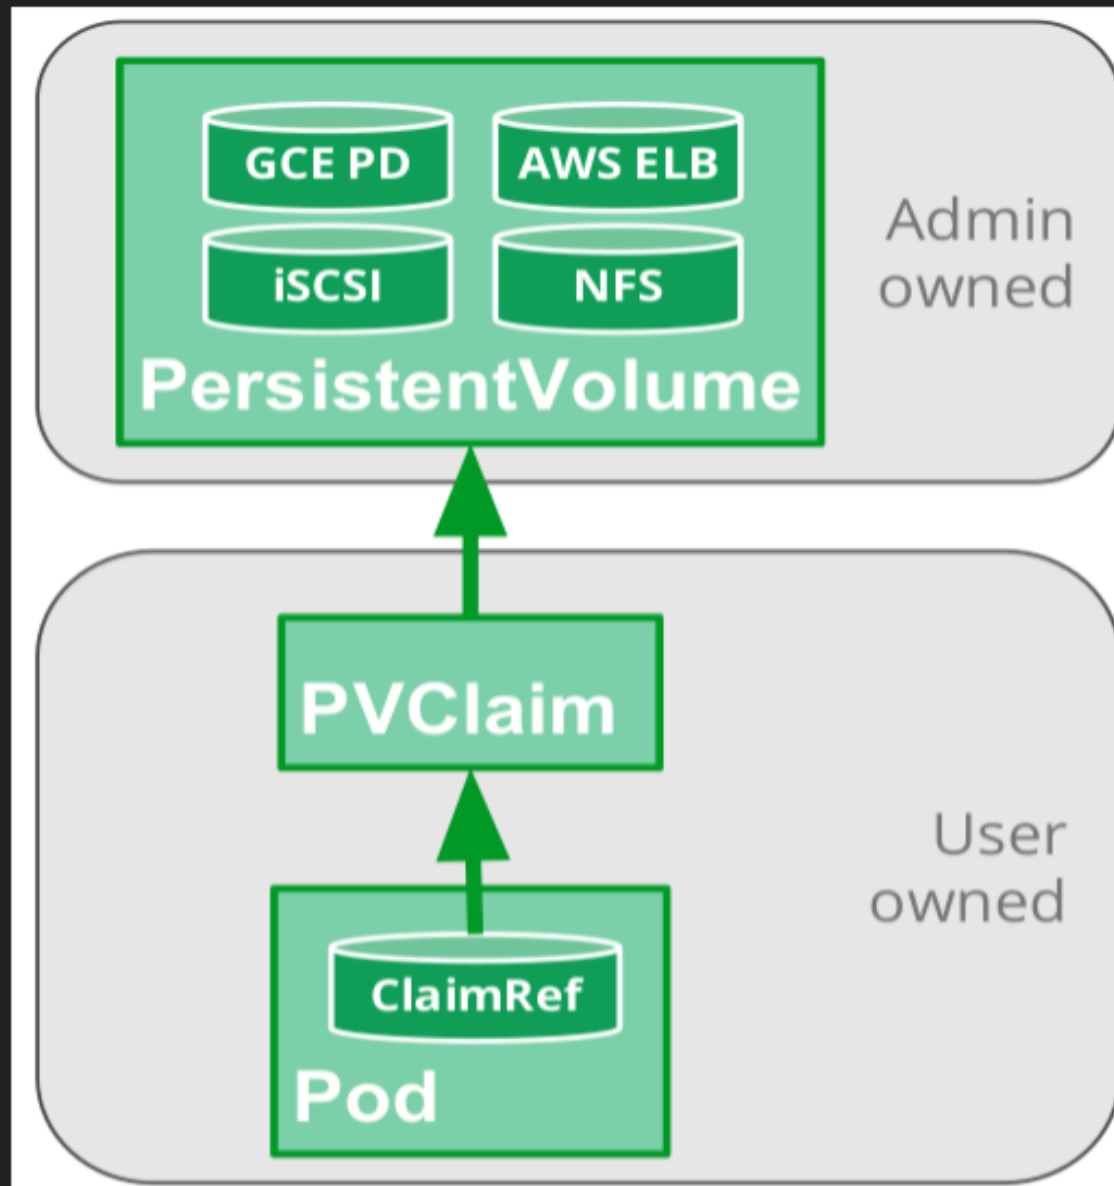
- Using volumes, you could deploy **applications with state** on your cluster
 - Those applications need to read/write to files on the **local filesystem** that need to be persistent in time
- You could run a MySQL database using persistent volumes
 - Although this might not be ready for production (yet)
 - Volumes are new since the June 2016 release in Kubernetes

Volumes

- If your **node stops** working, the pod can be rescheduled on another node, and the volumes can be attached to new node



Persistent Volumes



GCEPersistentDisk

CephFS

AWSElasticBlockStore

Cinder (OpenStack block storage)

AzureFile

Glusterfs

AzureDisk

VsphereVolume

FC (Fibre Channel)

Quobyte Volumes

FlexVolume

HostPath

Flocker

VMware Photon

NFS

vPortworx Volumes

iSCSI

ScaleIO Volumes

RBD (Ceph Block Device)

StorageOS

Persistent Volume Usage Design



Persistent Volume Provisioning Types

- **Static**
 - Manually map physical storage to PVs
- **Dynamic**
 - When static PVs are unavailable, K8S will check the availability of dynamic PVs
 - It dynamically creates volumes using storage provisioners
 - The storage provisioner may use a single physical storage for providing multiple dynamic PVs



Persistent Volume Claims

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```



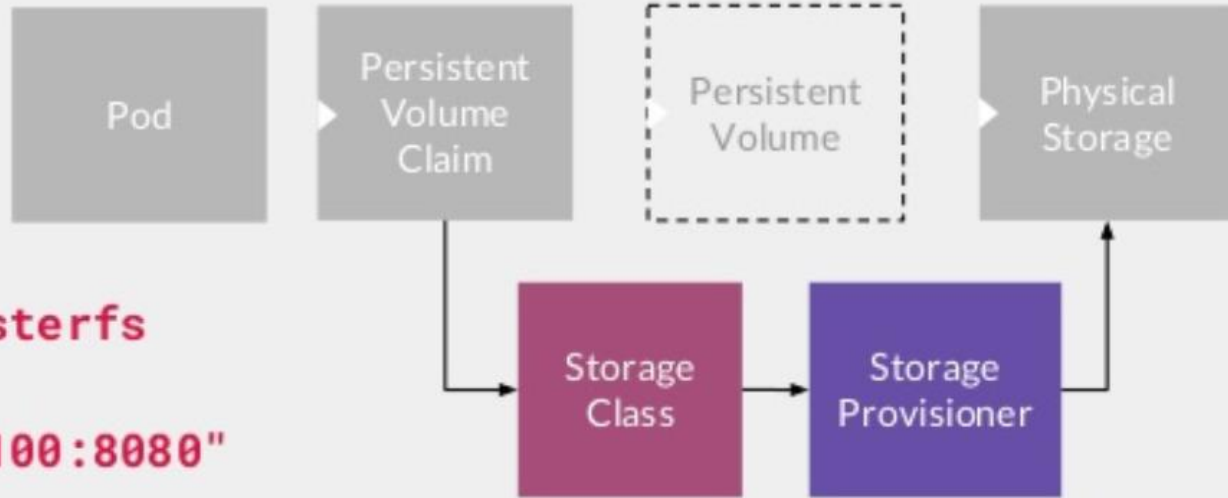
Static Persistent Volumes

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```



Dynamic Persistent Volumes

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: gluster-vol-default
provisioner: kubernetes.io/glusterfs
parameters:
  resturl: "http://192.168.10.100:8080"
  restuser: ""
  secretNamespace: ""
  secretName: ""
allowVolumeExpansion: true
```



Persistent Volume Volume Modes

- Filesystem
- Block

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: block-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  volumeMode: Block
  persistentVolumeReclaimPolicy: Retain
  fc:
    targetWWNs: ["50060e801049cfd1"]
    lun: 0
    readOnly: false
```



Persistent Volume Access Modes

- **ReadWriteOnce**
 - read-write by a single node
- **ReadOnlyMany**
 - read-only by many nodes
- **ReadWriteMany**
 - read-write by many nodes



Persistent Volume Binding

- Binding a PVC to PV will be done based on requested:
 - Amount of storage
 - Access mode
 - Storage Class (Optional)



Persistent Volume Re-claim Policies

- **Retain**
 - Manual reclamation
- **Recycle**
 - Basic scrub (`rm -rf /thevolume/*`) on the volume and makes it available again for a new claim
- **Delete**
 - Deletion removes both the PV object from Kubernetes, as well as the associated storage asset in the external infrastructure



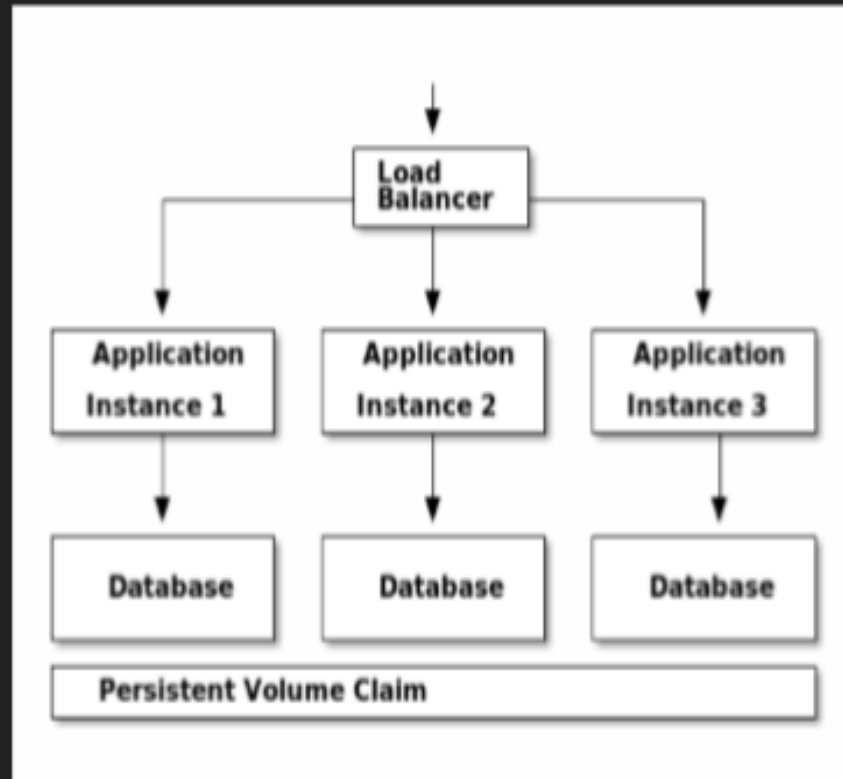
STORAGECLASS

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: standard
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
```

USING A STORAGECLASS

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 8Gi
  storageClassName: standard
```

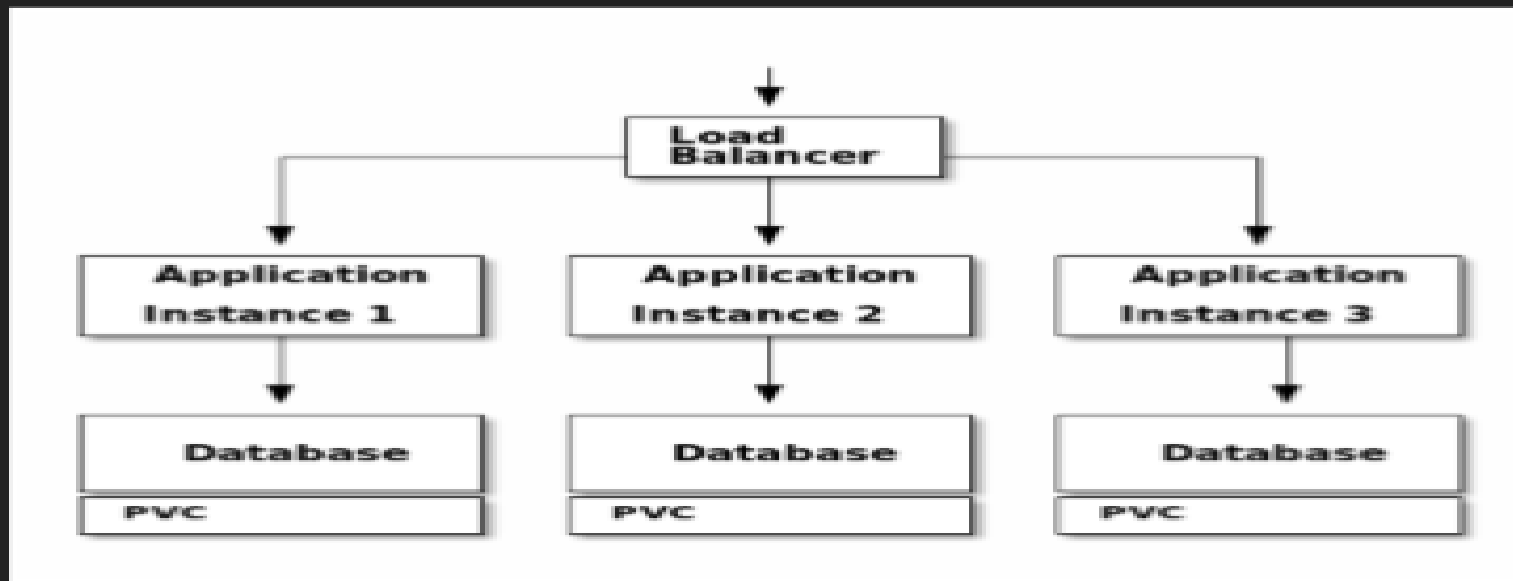
ARCHITECTURE REVISITED



SHARING PVCS

- Need different paths per Pod.
- Share needs to be available on every host.
- Corruption of file system affects all database nodes.
- Sharing a Persistent Volume Claim is not a good idea.

ARCHITECTURE STATEFULSETS



STATEFULSETS

- Pod template mechanism
- Hostnames are atomically increased:
 - pod-0
 - pod-1
 - ...
- Volume Claims can be provisioned on-the-fly

STATEFULSET EXAMPLE

```
apiVersion: apps/v1beta1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: gcr.io/google_containers/nginx-slim:0.8
          ports:
            - containerPort: 80
              name: web
          volumeMounts:
```