



Resource Quotas

- When a Kubernetes cluster is used by multiple **people** or **teams**, **resource management** becomes more important
 - You want to be able to **manage the resources** you give to a person or a team
 - You don't want one person or team **taking up all the resources** (e.g. CPU/Memory) of the cluster
- You can divide your cluster in **namespaces** (explained in next lecture) and enable resource quotas on it
 - You can do this using the **ResourceQuota** and **ObjectQuota** objects

Resource Quotas

- The administrator can set the following resource limits within a namespace:

Resource	Description
→ requests.cpu	The sum of CPU requests of all pods cannot exceed this value
requests.mem	The sum of MEM requests of all pods cannot exceed this value
requests.storage	The sum of storage requests of all persistent volume claims cannot exceed this value
limits.cpu	The sum of CPU limits of all pods cannot exceed this value
limits.memory	The sum of MEM limits of all pods cannot exceed this value

Resource Quotas

- The administrator can set the following object limits:

Resource	Description
configmaps	total number of configmaps that can exist in a namespace
persistentvolumeclaims	total number of persistent volume claims that can exist in a namespace
Pods	total number of Pods that can exist in a namespace
replicationcontrollers	total number of replicationcontrollers that can exist in a namespace
resourcequotas	total number of resource quotas that can exist in a namespace
services	total number of services that can exist in a namespace
services.loadbalancer	total number of load balancers that can exist in a namespace
services.nodeports	total number of nodeports that can exist in a namespace
secrets	total number of secrets that can exist in a namespace

Resource Quotas

- If a capacity quota (e.g. mem / cpu) has been specified by the administrator, then each pod needs to specify capacity quota during creation
 - The administrator can specify default request values for pods that don't specify any values for capacity
 - The same is valid for limit quotas
- If a resource is requested more than the allowed capacity, the server API will give an error 403 FORBIDDEN - and kubectl will show an error

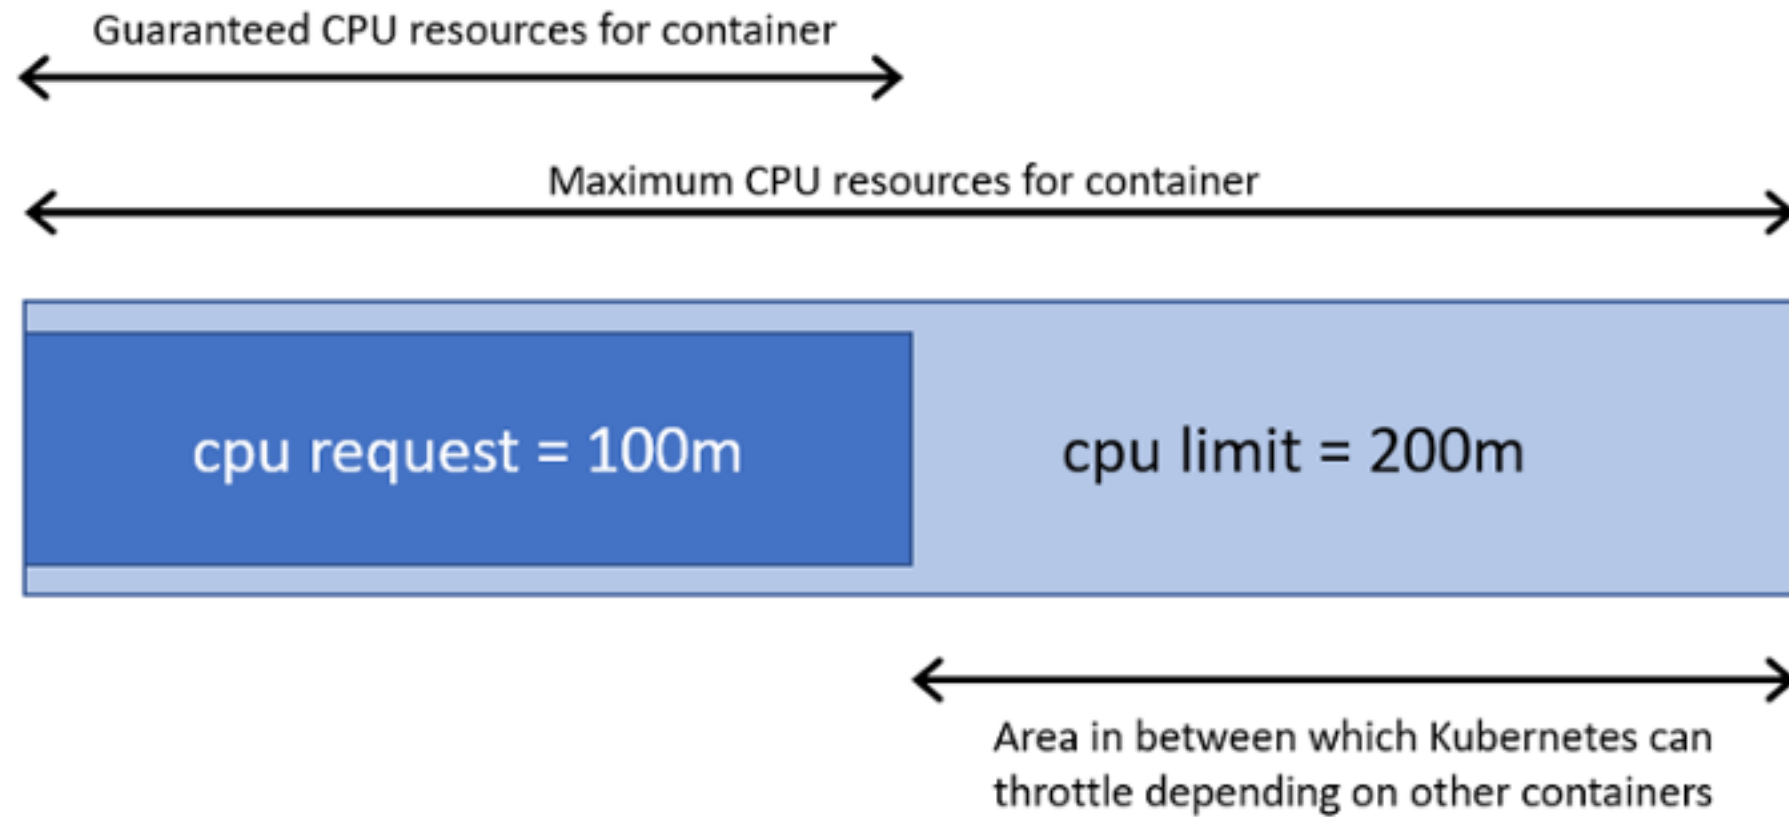


Resource Quotas

- Example of resource quotas:
 - You run a **deployment** with a **pod** with a **CPU resource** request of **200m**
 - 200m = 200 millicpu (or also 200 millicores)
 - 200m = 0.2, which is 20% of a CPU core of the running node
 - If the node has 2 cores, it's still 20% of a single core
 - You can also put a limit, e.g. on 400m
 - Memory quotas are defined by MiB or GiB



Quota and Limits



Resource Quotas

- Each container can specify **request capacity** and **capacity limits**
 - **Request capacity** is an **explicit request** for resources
 - The scheduler can use the **request capacity** to make decisions on where to put the pod on
 - You can see it as a **minimum amount of resources the pod needs**
 - **Resource limit** is a limit imposed to the container
 - The container will not be able to utilize more resources than specified

Namespaces

- You can then create resource limits within that namespace:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
  namespace: myspace
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```



Namespaces

- You can also create object limits:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: object-counts
  namespace: myspace
spec:
  hard:
    configmaps: "10"
    persistentvolumeclaims: "4"
    replicationcontrollers: "20"
    secrets: "10"
    services: "10"
    services.loadbalancers: "2"
```



